


Threads

- A thread is a basic unit of cpu utilization.
- It comprises of a thread ID, a program counter, a register set and a stack.
- A traditional process has a single thread of control. If a process has multiple thread of control. If a process has multiple thread of control it can perform more than one task at a time.

Multithreading

- Most modern applications are multithreaded
- Threads run within application
- Multiple tasks with the application can be implemented by separate threads
 - Update display
 - Fetch data
 - Spell checking
 - Answer a network request
- Process creation is heavy-weight while thread creation is light-weight
- Can simplify code, increase efficiency
- Kernels are generally multithreaded

Multithreaded Server Architecture



Benefits

- **Responsiveness** – May allow continued execution if part of process is blocked, especially important for user interfaces.
For instance a web browser could allow user interaction in one thread while an image was loaded in another thread.
- **Resource Sharing** – Threads share resources of process, easier than shared memory or message passing. *Benefit of sharing code and data allows application to have several threads of activity within the same address space.*
- **Economy** – Cheaper than process creation, thread switching lower overhead than context switching. In general it is more time consuming to create and manage process than threads.
- **Scalability** – Process can take advantage of multiprocessor architectures. Single threaded process can run on one processor. Multi threading on a multi cpu machine increases parallelism.
- *For Ex. A word processor(e.g., MS word) may have a thread for displaying graphics, another thread for keystrokes, another thread for spell checking and so on.*

Multicore Programming


- **Multicore** or **multiprocessor** systems putting pressure on programmers, challenges include:
 - **Dividing activities:** find areas that can be divided into separate concurrent task.
 - **Balance:** balance of task equal work as task run in parallel
 - **Data splitting:** Data accessed and manipulated must be divided to run in separate areas.
 - **Data dependency:** one task depends on other tasks data therefore execution of task must be synchronized.
 - **Testing and debugging:** Different testing and debugging task for programs must run in parallel on multiple cores
- **Parallelism** implies a system can perform more than one task simultaneously
- **Concurrency** supports more than one task making progress
 - Single processor / core, scheduler providing concurrency

Multicore Programming (Cont.)


- Types of parallelism
 - **Data parallelism** - distributes subsets of the same data across multiple cores, same operation on each
 - **Task parallelism** - distributing threads across cores, each thread performing unique operation
- As # of threads grows, so does architectural support for threading
 - CPUs have cores as well as *hardware threads*
 - Consider Oracle SPARC T4 with 8 cores, and 8 hardware threads per core

Concurrency vs. Parallelism

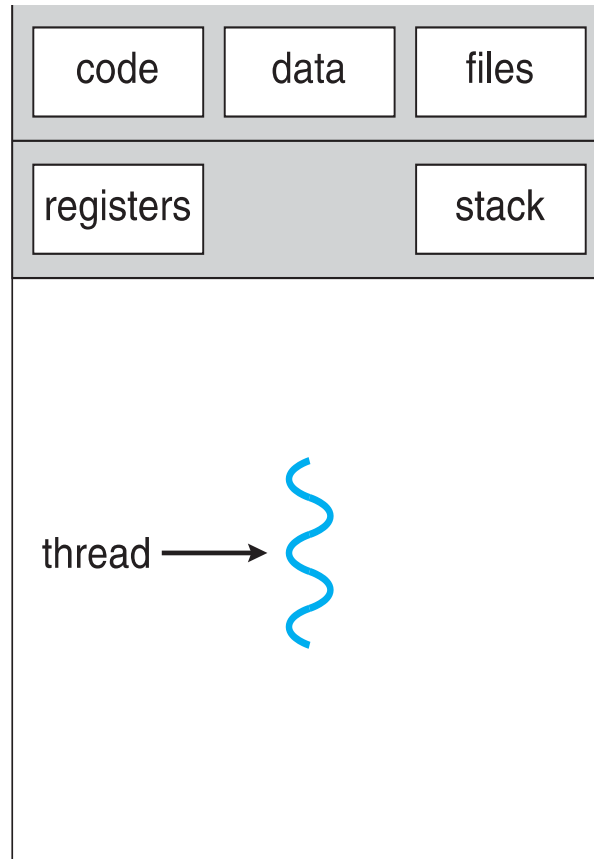
■ Concurrent execution on single-core system:



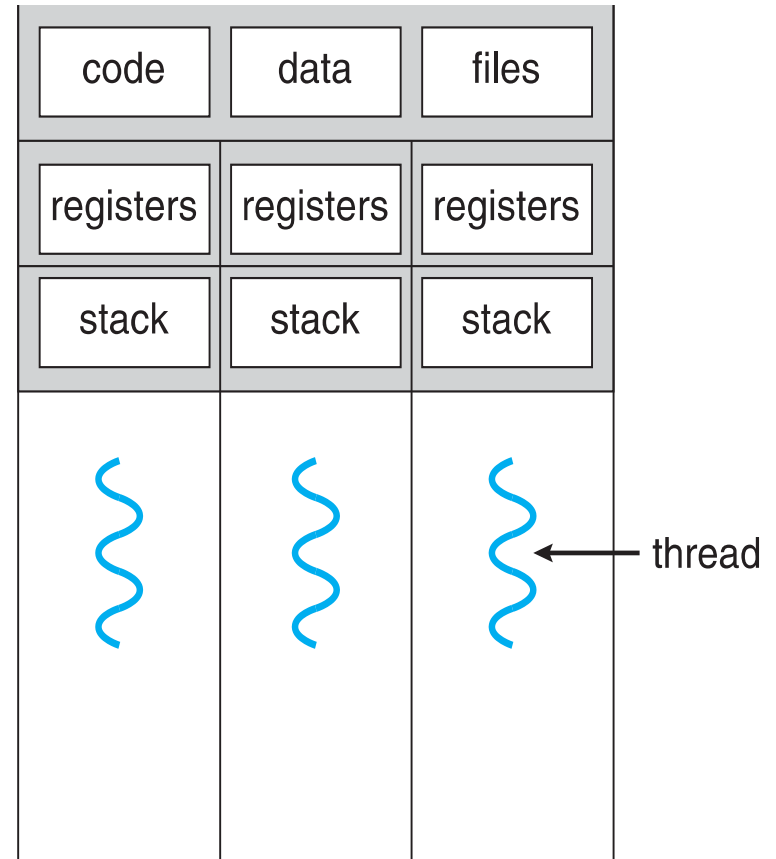
■ Parallelism on a multi-core system:



Single and Multithreaded Processes



single-threaded process



multithreaded process

Amdahl's Law

- Amdahl's Law is a formulation, an observation that attempts to give an upper bound on the performance increase that can be realized when a fixed workload is moved to a more parallel environment (e.g. multiprocessing or multi-core environment).
- S is serial portion
- N processing cores

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$

- That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
- As N approaches infinity, speedup approaches $1 / S$

Serial portion of an application has disproportionate effect on performance gained by adding additional cores

Amdahl's Law

- You can also consider $\text{speedup} \leq \frac{1}{(1-p+\frac{p}{N})}$ where p = parallel portion
and N = processing cores
- If 30% of the execution time may be the subject of a speedup, S will be 0.7 and $p=0.3$; if the improvement makes the affected part twice as fast, N will be 2. Amdahl's law states that the overall speedup of applying the improvement will be????
- Assume that we are given a serial task which is split into four consecutive parts, whose percentages of execution time are $p_1 = 0.11$, $p_2 = 0.18$, $p_3 = 0.23$, and $p_4 = 0.48$ respectively. Then we are told that the 1st part is not sped up, so $N_1 = 1$, while the 2nd part is sped up 5 times, so $N_2 = 5$, the 3rd part is sped up 20 times, so $N_3 = 20$, and the 4th part is sped up 1.6 times, so $N_4 = 1.6$. By using Amdahl's law, Find the overall speedup ????

User Threads and Kernel Threads


- **User threads** - management done by user-level threads library
- Three primary thread libraries:
 - POSIX **Pthreads**
 - Windows threads
 - Java threads
- **Kernel threads** - Supported by the Kernel
- Examples - virtually all general purpose operating systems, including:
 - Windows
 - Solaris
 - Linux
 - Tru64 UNIX
 - Mac OS X

Multithreading Models

- *Many-to-One*
- *One-to-One*
- *Many-to-Many*


Many-to-One

- Many user-level threads mapped to single kernel thread
- One thread blocking causes all to block
- Multiple threads may not run in parallel on muticore system because only one may be in kernel at a time
- Few systems currently use this model
- Examples:
 - Solaris Green Threads
 - GNU Portable Threads




One-to-One

- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead
- Examples
 - Windows
 - Linux
 - Solaris 9 and later




Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Solaris prior to version 9
- Windows with the *ThreadFiber* package



Two-level Model

- Similar to M:M, except that it allows a user thread to be **bound** to kernel thread
- Examples
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier



End of Lecture

Thank You