



OPERATING SYSTEMS

CSE-4041

Lecture-6

CPU Scheduling

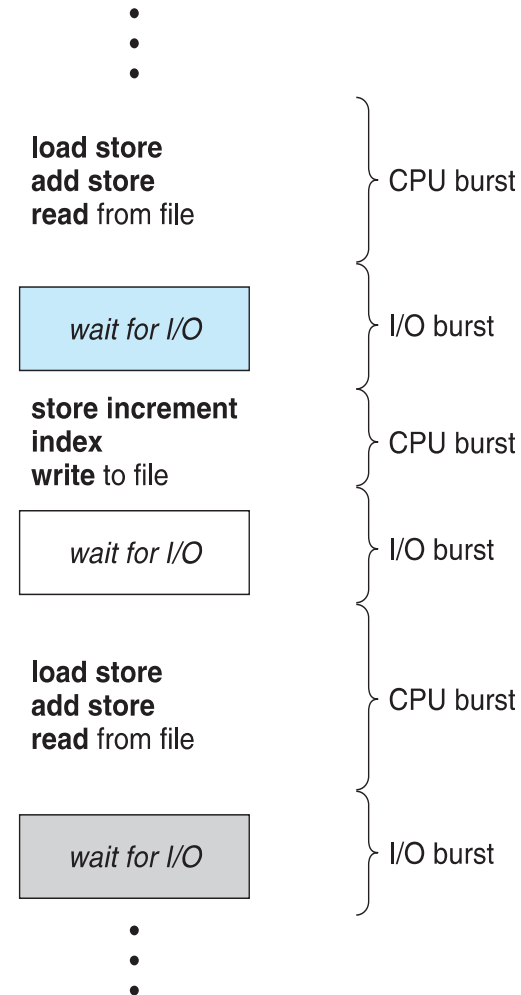
Nitesh Kumar Jha

Assistant Professor

ITER S'O'A (Deemed to be University)

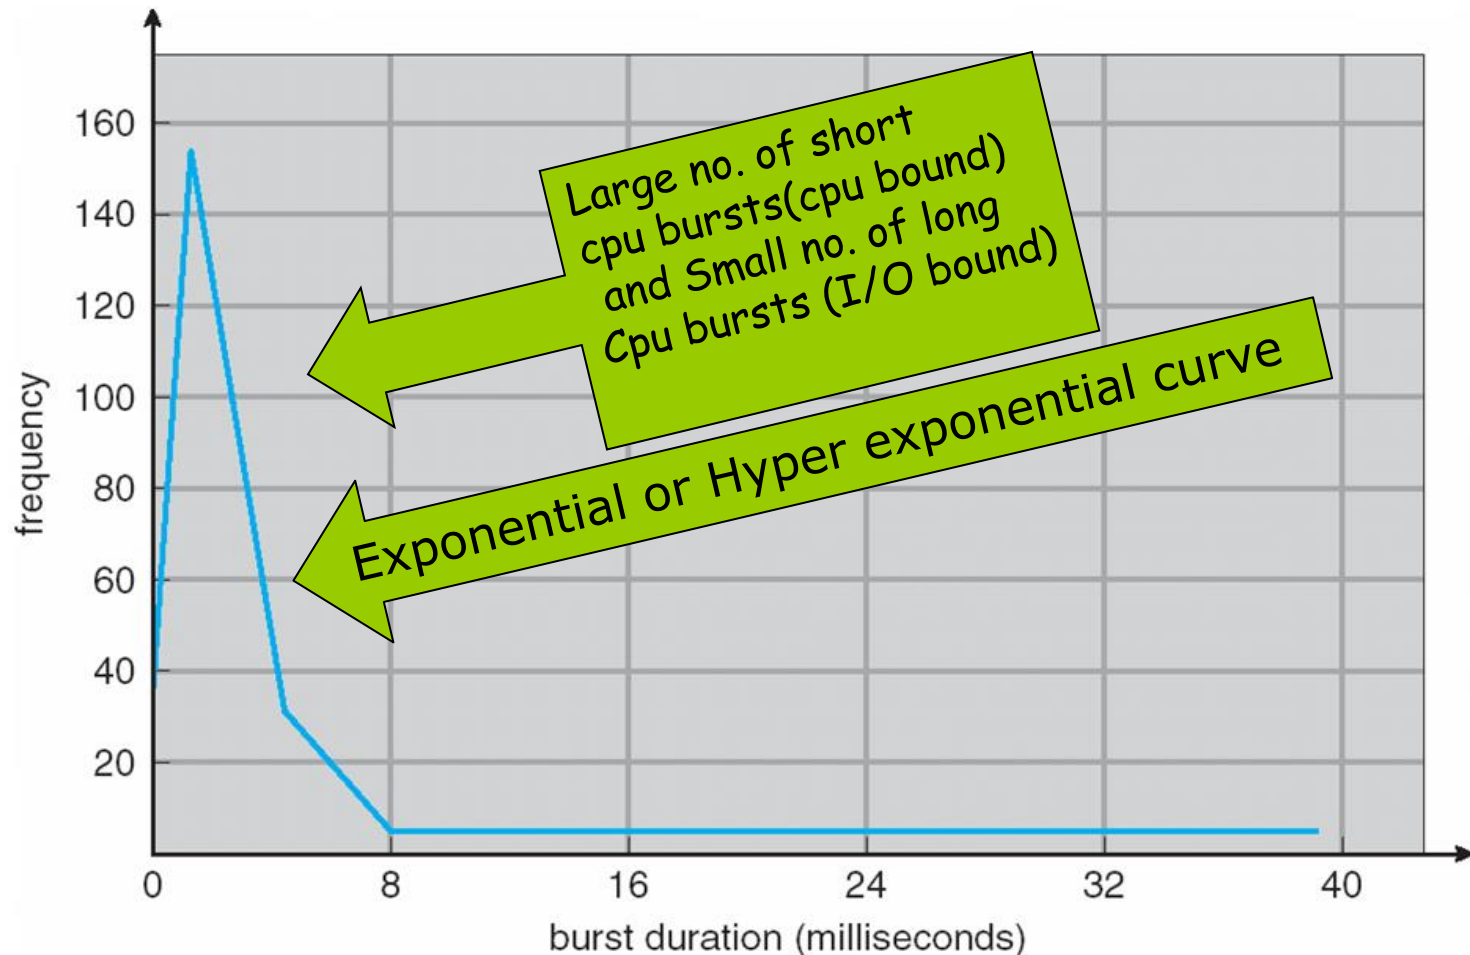
CPU SCHEDULING

- Maximum CPU utilization obtained with multiprogramming
- CPU-I/O Burst Cycle - Process execution consists of a **cycle** of CPU execution and I/O wait
- **CPU burst** followed by **I/O burst**
- CPU burst distribution is of main concern



ALTERNATE sequence for CPU and I/O bursts

Histogram of CPU-burst Times



Histogram of CPU-burst durations

CPU Scheduler

- **Short-term scheduler** selects from among the processes in ready queue, and allocates the CPU to one of them
 - Queue may be ordered in various ways
- CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state(I/O request or execution of fork())
 2. Switches from running to ready state(Interrupt occurs)
 3. Switches from waiting to ready(completion of I/O)
 4. Process Terminates
- Scheduling under 1 and 4 is **non preemptive**
- All other scheduling is **preemptive**
 - Consider access to shared data
 - Consider preemption while in kernel mode
 - Consider interrupts occurring during crucial OS activities

Dispatcher

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - switching context
 - switching to user mode
 - jumping to the proper location in the user program to restart that program
- **Dispatch latency** - time it takes for the dispatcher to stop one process and start another running

Scheduling Criteria

- **CPU utilization** - keep the CPU as busy as possible, 40% lightly loaded-90% heavily loaded.
- **Throughput** - # of processes that complete their execution per time unit. For long process it may be 1 process per hour or may be 10 processes per second(short processes).
- **Turnaround time** - amount of time to execute a particular process i.e the time starting from the submission of process upto its termination.
 - $TAT = \text{completion time} - \text{Arrival time}$
- **Waiting time** - amount of time a process has been waiting in the ready queue.
 - $WT = TAT - \text{Burst time}$
- **Response time** - amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

Scheduling Algorithm Optimization Criteria

Ideally for any CPU scheduling Algorithm it must satisfy below:

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time

CPU scheduling Algorithms:

It is the task of selecting any one process from the ready queue and allocating the cpu for execution.

First- Come, First-Served (FCFS) Scheduling

- Simplest of all scheduling algorithms.
- First process to enter the ready queue shall be allocated the CPU.
- This algorithm is non preemptive and doesn't suit time sharing system.

Consider three process P_1, P_2 and P_3 with their respective burst times

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

What are the possible order in which processes arrive?

First- Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Suppose that the processes arrive in the order: P_1, P_2, P_3
The Gantt Chart for the schedule is:



- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$
- Turn around time $P_1 = 24$; $P_2 = 27$; $P_3 = 30$
- Average turn around time: $(24+27+30)/3=27$

FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

- The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- **Convoy effect** - short process behind long process
 - Consider one CPU-bound and many I/O-bound processes


Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst
 - Use these lengths to schedule the process with the shortest time
- SJF is optimal - gives minimum average waiting time for a given set of processes
 - The difficulty is knowing the length of the next CPU request
 - Could ask the user

Example of SJF

<u>Process</u>	<u>Burst Time</u>
P_1	6
P_2	8
P_3	7
P_4	3

■ SJF scheduling chart



■ Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$


Implementation of SJF

- This Algorithm is optimal in terms of average waiting time for a set of processes.
- It is not suitable for CPU scheduling cause it is impossible to calculate or predict the next CPU time of the process waiting in the ready queue.
- The SJF algorithm is somehow used for long term scheduling performed by the job scheduler where the burst time of a process can be calculated.
- Although SJF is impossible one approach of approximation can be done by using the length of previous CPU bursts, using exponential averaging

1. t_n = actual length of n^{th} CPU burst
2. τ_{n+1} = predicted value for the next CPU burst
3. $\alpha, 0 \leq \alpha \leq 1$
4. Define : $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$.

- Commonly, α is set to $\frac{1}{2}$
- Preemptive version called **shortest-remaining-time-first(SRTF)**

Prediction of the Length of the Next CPU Burst



CPU burst (t_i)	6	4	6	4	13	13	13	...
"guess" (τ_i)	10	8	6	6	9	11	12	...

Examples of Exponential Averaging

■ $\alpha = 0$

- $\tau_{n+1} = \tau_n$
- Recent history does not count

■ $\alpha = 1$

- $\tau_{n+1} = \alpha t_n$
- Only the actual last CPU burst counts

■ If we expand the formula, we get:

$$\begin{aligned}\tau_{n+1} = & \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ & + (1 - \alpha)^{n+1} \tau_0\end{aligned}$$


- ## ■ Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor

Example of Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

- Preemptive SJF Gantt Chart*



- Average waiting time = $[(10-1)+(1-1)+(17-2)+5-3]/4 = 26/4 = 6.5$ milli seconds

Priority Scheduling

- A priority number (integer) is associated with each process, The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
- Priority assigned to a process may either be
 - internally defined (by the OS)
 - ▶ ratio of CPU to I/O burst
 - ▶ Overall memory requirements
 - Externally defined
 - ▶ Based on importance of process
 - ▶ Funding or sponsorship or other administrative factors
- SJF is priority scheduling where priority is the inverse of predicted next CPU burst time


Priority Scheduling

- Priority scheduling can either be
 - Preemptive
 - Non preemptive
- In cases where two or more processes are found to have same priority then FCFS logic is used to break the tie.
- Problem \equiv **Starvation** - In certain cases low priority processes may never execute as there is a flow of higher priority processes.
- Solution \equiv **Aging** - as time progresses increase the priority of the process.

Priority Scheduling(Non Preemptive)

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

■ Priority scheduling Gantt Chart




■ Average waiting time = 8.2 m.sec

Priority Scheduling(Preemptive)

<u>Process</u>	<u>Arrival Time</u>	<u>Priority</u>	<u>Burst time</u>
P_1	0	3	10
P_2	1	1	1
P_3	2	4	2
P_4	4	2	3

■ Priority scheduling Gantt Chart



■ Average waiting time = 4 m.sec


Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum q**), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- Timer interrupts every quantum to schedule next process
- Performance
 - q large \Rightarrow FIFO
 - q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high

Example of RR with Time Quantum = 4


<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- The Gantt chart is:




- Typically, higher average turnaround than SJF, but better response
- q should be large compared to context switch time
- q usually 10ms to 100ms, context switch < 10 usec

Time Quantum and Context Switch Time



Turnaround Time Varies With The Time Quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7


80% of CPU bursts
should be shorter than q

Multilevel Queue

- Ready queue is partitioned into separate queues, eg:
 - **foreground** (interactive)
 - **background** (batch)
- Process permanently in a given queue
- Each queue has its own scheduling algorithm:
 - foreground - RR
 - background - FCFS
- Scheduling must be done between the queues:
 - Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
 - Time slice - each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR
 - 20% to background in FCFS

Multilevel Queue Scheduling

highest priority



lowest priority

Multilevel Feedback Queue

- A process can move between the various queues; aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service

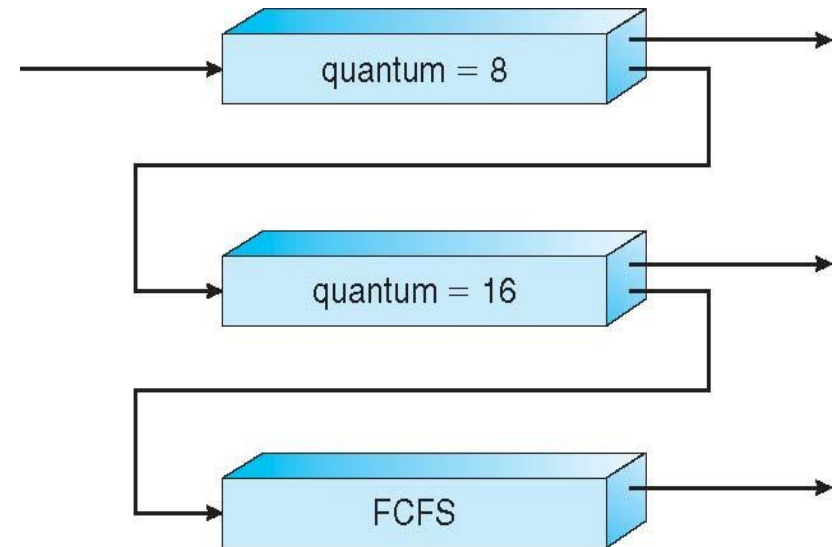
Example of Multilevel Feedback Queue

■ Three queues:

- Q_0 - RR with time quantum 8 milliseconds
- Q_1 - RR time quantum 16 milliseconds
- Q_2 - FCFS

■ Scheduling

- A new job enters queue Q_0 which is served FCFS
 - ▶ When it gains CPU, job receives 8 milliseconds
 - ▶ If it does not finish in 8 milliseconds, job is moved to queue Q_1
- At Q_1 job is again served FCFS and receives 16 additional milliseconds
 - ▶ If it still does not complete, it is preempted and moved to queue Q_2




Thread Scheduling

- Distinction between user-level and kernel-level threads
- When threads supported, threads scheduled, not processes
- Many-to-one and many-to-many models, thread library schedules user-level threads to run on LWP
 - Known as **process-contention scope (PCS)** since scheduling competition is within the process
 - Typically done via priority set by programmer
- Kernel thread scheduled onto available CPU is **system-contention scope (SCS)** - competition among all threads in system

Real-Time CPU Scheduling

- Can present obvious challenges
- **Soft real-time systems** - no guarantee as to when critical real-time process will be scheduled
- **Hard real-time systems** - task must be serviced by its deadline
- Two types of latencies affect performance
 1. Interrupt latency - time from arrival of interrupt to start of routine that services interrupt
 2. Dispatch latency - time for schedule to take current process off CPU and switch to another



End of Lecture

Thank You