

Information Retrieval

Topic- Index Construction (Hardware basics, BSBI)

Lecture-17

Prepared By

Dr. Rasmita Rautray & Dr. Rasmita Dash

Associate Professor

Dept. of CSE

Content

- Hardware basics
- Block Sort Based Indexing (BSBI)

Hardware basics

- Many design decisions in information retrieval are based on hardware constraints.
- We begin by reviewing hardware basics that we'll need in this course.

Hardware basics

- Access to data is much faster in memory than on disk. (roughly a factor of 10)
- Disk seeks are “idle” time: No data is transferred from disk while the disk head is being positioned.
- To optimize transfer time from disk to memory: one large chunk is faster than many small chunks.
- Disk I/O is block-based: Reading and writing of entire blocks (as opposed to smaller chunks). Block sizes: 8KB to 256 KB
- Servers used in IR systems typically have many GBs of main memory and TBs of disk space.
- Fault tolerance is expensive: It's cheaper to use many regular machines than one fault tolerant machine.

Data collection


- Shakespeare's collected works are not large enough for demonstrating many of the points in this course.
- As an example, Reuters RCV1 (Reuters Corpus Volume 1) collection will be used for applying scalable index construction algorithms.
 - It is an English newswire articles sent over the wire in 1995 and 1996 (one year).

RCV1 statistics

N : documents	800,000
L : tokens per document	200
M : terms (= word types)	400,000
bytes per token (incl. spaces/punct.)	6
bytes per token (without spaces/punct.)	4.5
bytes per term (= word type)	7.5
T : non-positional postings	100,000,000

Index construction

Goal: construct the inverted index



Sort-based index construction

- As we build index, we parse docs one at a time.
- The final postings for any term are incomplete until the end.
- Can we keep all postings in memory and then do the sort in-memory at the end?
- No, not for large collections
- Thus: We need to store intermediate results on disk.

Same algorithm for disk?

- Can we use the same index construction algorithm for larger collections, but by using disk instead of memory?
- No: Sorting very large sets of records on disk is too slow – too many disk seeks.
- We need an external sorting algorithm.

“External” sorting algorithm (using few disk seeks)

- We must sort $T = 100,000,000$ non-positional postings.
 - Each posting has size 12 bytes (4+4+4: termID, docID, term frequency).
- Define a block to consist of 10,000,000 such postings
 - We can easily fit that many postings into memory.
 - We will have 10 such blocks for RCV1.
- Basic idea of algorithm:
 - For each block: (i) accumulate postings, (ii) sort in memory, (iii) write to disk
 - Then merge the blocks into one long sorted order.

Merging two blocks

postings lists
to be merged


brutus	d1,d3
caesar	d1,d2,d4
noble	d5
with	d1,d2,d3,d5

brutus	d6,d7
caesar	d8,d9
julius	d10
killed	d8



brutus	d1,d3,d6,d7
caesar	d1,d2,d4,d8,d9
julius	d10
killed	d8
noble	d5
with	d1,d2,d3,d5

merged
postings lists



Blocked Sort-Based Indexing

BSBINDEXCONSTRUCTION()

```
1   $n \leftarrow 0$   
2  while (all documents have not been processed)  
3  do  $n \leftarrow n + 1$   
4      $block \leftarrow \text{PARSENEXTBLOCK}()$   
5     BSBI-INVERT( $block$ )  
6     WRITEBLOCKTODISK( $block, f_n$ )  
7  MERGEBLOCKS( $f_1, \dots, f_n; f_{\text{merged}}$ )
```

Thank You