



# Transaction Management

Instructor : Nitesh Kumar Jha

[niteshjha@soa.ac.in](mailto:niteshjha@soa.ac.in)

ITER,S'O'A(DEEMED TO BE  
UNIVERSITY)

Sept 2018

# Transaction Concept


- A **transaction** is a *unit* of program execution that accesses and possibly updates various data items.
- E.g., transaction to transfer Rs. 50 from account A to account B:

1. **read**(A)
2.  $A := A - 50$
3. **write**(A)
4. **read**(B)
5.  $B := B + 50$
6. **write**(B)

One Transaction consists of a Set of instruction

- Two main issues to deal with:
  - Failures of various kinds, such as hardware failures and system crashes
  - Concurrent execution of multiple transactions

# Transaction State - I



- **Active** – the initial state; the transaction stays in this state while it is executing
- **Partially committed** – after the final statement has been executed.
- **Failed** .. after the discovery that normal execution can no longer proceed.

- **Aborted** – after the transaction has been rolled back and the database restored to its state prior to the start of the transaction. Two options after it has been aborted:
  - Restart the transaction
    - ▶ can be done only if no internal logical error
  - Kill the transaction
- **Committed** – after successful completion.

# ACID Properties

To preserve the integrity of data the database system must ensure:

- **Atomicity.** Either all operations of the transaction are properly reflected in the database or none are.
- **Consistency.** Execution of a transaction in isolation preserves the consistency of the database.
- **Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.
  - That is, for every pair of transactions  $T_i$  and  $T_j$ , it appears to  $T_i$  that either  $T_j$  finished execution before  $T_i$  started, or  $T_j$  started execution after  $T_i$  finished.
- **Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.

# ACID Properties w.r.t. a Transaction

## ■ Atomicity requirement

- If the transaction fails after step 3, money will be “lost” leading to an inconsistent database state
- The system should ensure that updates of a partially executed transaction are not reflected in the database

1. **read(A)**
2.  $A := A - 50$
3. **write(A)**
4. **read(B)**
5.  $B := B + 50$
6. **write(B)**

- **Durability requirement** —The database should be durable enough to hold all its latest updates even if the system fails or restarts. If a transaction updates a chunk of data in a database and commits, then the database will hold the modified data. If a transaction commits but the system fails before the data could be written on to the disk, then that data will be updated once the system springs back into action.

## ■ Consistency requirement in above example:

- The sum of A and B is unchanged by the execution of the transaction

# Required Properties of a Transaction (Cont.)

- **Isolation requirement** — if between steps 3 and 6, another transaction **T2** is allowed to access the partially updated database

**T1**

1. **read**(A)
2.  $A := A - 50$
3. **write**(A)
4. **read**(B)
5.  $B := B + 50$
6. **write**(B)

**T2**

read(A), read(B), print(A+B)

- Isolation can be ensured trivially by running transactions **serially**
  - That is, one after the other.
- However, executing multiple transactions concurrently has significant benefits.

# Concurrent Executions

- Multiple transactions are allowed to run concurrently in the system. Advantages are:
  - **Increased processor and disk utilization**, leading to better transaction *throughput*
    - ▶ E.g. one transaction can be using the CPU while another is reading from or writing to the disk
  - **Reduced average response time** for transactions: short transactions need not wait behind long ones.
- **Concurrency control schemes** – mechanisms to achieve isolation
  - That is, to control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database

# Schedules

- **Schedule** – a sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed
  - A schedule for a set of transactions must consist of all instructions of those transactions
  - Must preserve the order in which the instructions appear in each individual transaction.
- A transaction that successfully completes its execution will have a **commit** instructions as the last statement
  - By default transaction assumed to execute commit instruction as its last step
- A transaction that fails to successfully complete its execution will have an **abort** instruction as the last statement



# Serial Schedule 1

- Let  $T_1$  transfer Rs.50 from A to B, and  $T_2$  transfer 10% of the balance from A to B.
- An example of a **serial** schedule in which  $T_1$  is followed by  $T_2$  :

$T_1$	$T_2$
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	

Serial execution of transactions  
always ensures isolation and  
consistence in database

read (A)  
 $temp := A * 0.1$   
 $A := A - temp$   
write (A)  
read (B)  
 $B := B + temp$   
write (B)  
commit

Consistency is preserved  
i.e.  $A+B$  remains same

# Serial Schedule 2

- A **serial** schedule in which  $T_2$  is followed by  $T_1$  :

$T_1$	$T_2$
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Consistency is preserved  
irrespective of execution sequence  
of both  $T_1$  and  $T_2$

# Concurrent Schedule 3

$T_1$	$T_2$
read (A) $A := A - 50$	
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B)
write (A) read (B) $B := B + 50$ write (B) commit	
	$B := B + temp$ write (B) commit

- concurrent schedule does not preserve the sum of A+B

# Concurrent Schedule 4

$T_1$	$T_2$
read (A) $A := A - 50$ write (A)	
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	
	read (B) $B := B + temp$ write (B) commit

- the sum “A + B” is preserved
- It is not a serial schedule, but it is **equivalent** to a serial Schedule.
- These schedules are called **serializable** schedules.
- i.e. out of multiple possible concurrent schedules, some may ensure isolation and other may not.
- Hence only the concurrent schedules that ensures isolation and consistency shall be acceptable.

# Serializability

- If the final outcome of a concurrent schedule  $S_1$ , is same as that of a serial schedule  $S_2$ , then  $S_1$  is said to be a serializable schedule.
- i.e. A concurrent schedule is serializable if it is equivalent to a serial schedule. Different forms of schedule equivalence give rise to the notions of:
  1. **Conflict Serializability**
  2. **View Serializability**
- Simplified view of transaction
  - We ignore operations other than **read** and **write** instructions
  - We assume that transactions may perform arbitrary computations in between reads and writes

# Conflicting Instructions

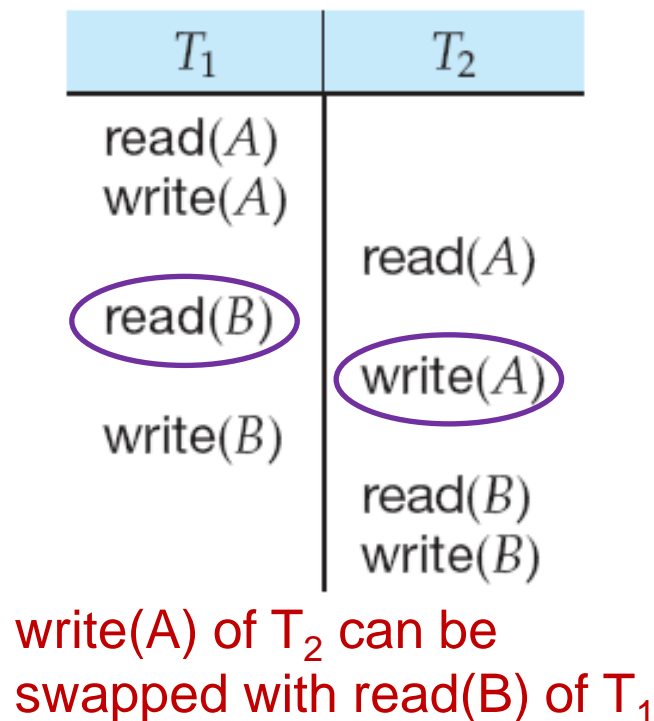
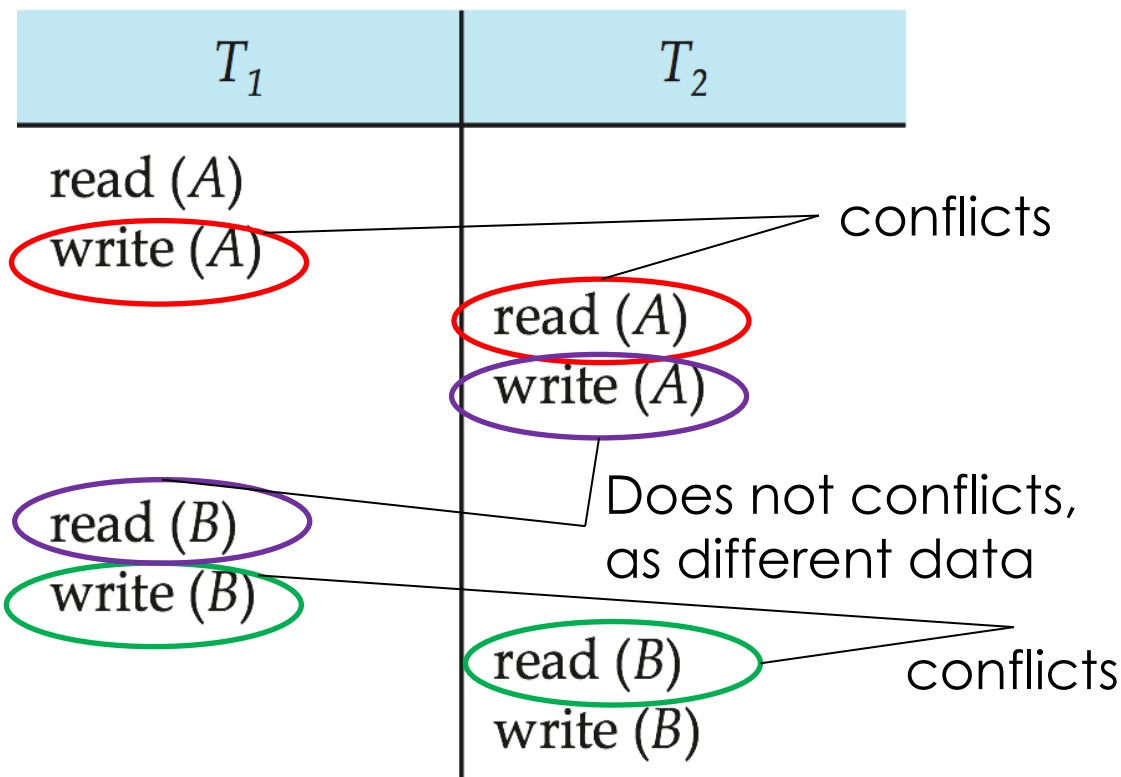
- Let  $I_i$  and  $I_j$  be two Instructions of transactions  $T_i$  and  $T_j$  respectively. Instructions  $I_i$  and  $I_j$  **conflict** if and only if there exists some item  $Q$  accessed by both  $I_i$  and  $I_j$ , and at least one of these instructions wrote  $Q$ .
  - 1.  $I_i = \text{read}(Q)$ ,  $I_j = \text{read}(Q)$ .  $I_i$  and  $I_j$  don't conflict, order does not matter
  - 2.  $I_i = \text{read}(Q)$ ,  $I_j = \text{write}(Q)$ . They conflict, as the order matters
  - 3.  $I_i = \text{write}(Q)$ ,  $I_j = \text{read}(Q)$ . They conflict, as the order matters
  - 4.  $I_i = \text{write}(Q)$ ,  $I_j = \text{write}(Q)$ . They conflict, order does not affect.

However, the value obtained by the next  $\text{read}(Q)$  is affected, since the result of only the latter of the two write instructions is preserved in the database

- Intuitively, a conflict between  $I_i$  and  $I_j$  forces a (logical) temporal order between them.
  - If  $I_i$  and  $I_j$  are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.
  - i.e. if both  $I_i$  and  $I_j$  represent read operation, then they can be swapped, but if any one of them is a write operation then they can not be swapped.

# Conflict Serializability

- If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions, we say that  $S$  and  $S'$  are **conflict equivalent**.
- We say that a schedule  $S$  is **conflict serializable** if it is conflict equivalent to a serial schedule



# Conflict Serializability (Cont.)

$T_1$	$T_2$	$T_1$	$T_2$
read (A) write (A)		read (A) write (A)	
	read (A) write (A)	read(B) write(B)	
read (B) write (B)			read(A) write(A) read (B) write (B)
	read (B) write (B)		

- Swap the read(B) instruction of  $T_1$  with the read(A) instruction of  $T_2$ .
- Swap the write(B) instruction of  $T_1$  with the write(A) instruction of  $T_2$ .
- The final result of these swaps is a serial schedule
- i.e. S and S' are **conflict equivalent**
- and hence S is **conflict serializable**



# Conflict Serializability (Cont.)

- Example of a schedule that is not conflict serializable:

$T_3$	$T_4$
read (Q)	
write (Q)	write (Q)


- It is not possible to swap instructions in the above schedule to obtain a serial schedule

# Precedence Graph

- It is a simple and efficient method for determining conflict serializability
- Consider schedule  $S$  of a set of transactions  $T_1, T_2, \dots, T_n$
- **Precedence graph** — a direct graph  $G=(V,E)$ ,
  - where the vertices are participating transactions.
  - We draw a directed from  $T_i$  to  $T_j$  if the two transaction conflict, i.e.
    - $T_i$  executes **write(Q)** before  $T_j$  executes **read(Q)**
    - $T_i$  executes **read(Q)** before  $T_j$  executes **write(Q)**
    - $T_i$  executes **write(Q)** before  $T_j$  executes **write(Q)**
- If the precedence graph for  $S$  has a **cycle**, then  $S$  is **not conflict serializable**

# Precedence graph for Schedule 1

$T_1$	$T_2$
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit



- Since all instructions of  $T_1$  are executed before the first instruction of  $T_2$  is executed.
- An edge is formed from  $T_1$  to  $T_2$
- As there is no cycle, therefore  $S_1$  is conflict serializable

# Concurrent Schedule 4


$T_1$	$T_2$
read (A) $A := A - 50$ write (A)	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	read (B) $B := B + temp$ write (B) commit



- $T_1$  executes write(A) before  $T_2$  executes read(A)
- $T_1$  executes read(B) before  $T_2$  executes write(B)
- $T_1$  executes write(B) before  $T_2$  executes write(B)
- As there is no cycle, therefore  $S_4$  is conflict serializable

# Precedence Graph Schedule 3


$T_1$	$T_2$
<code>read (A)</code> <code>A := A - 50</code>	<code>read (A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write (A)</code> <code>read (B)</code>
<code>write (A)</code> <code>read (B)</code> <code>B := B + 50</code> <code>write (B)</code> <code>commit</code>	<code>B := B + temp</code> <code>write (B)</code> <code>commit</code>



- One edge from  $T_1 \rightarrow T_2$ , as  $T_1$  executes `read(A)`, before  $T_2$  executes `write(A)`
- Another edge  $T_2 \rightarrow T_1$ , as  $T_2$  executes `read(B)`, before  $T_1$  executes `write(B)`
- As the precedence graph contains a **cycle**, therefore  $S_3$  is **not conflict serializable**

# Concurrent Schedule 5

$T_1$	$T_2$
Read(A) $A = A - 50$ Write(A)	Read(B) $B = B - 10$ Write(B)
Read(B) $B = B + 50$ Write(B)	Read(A) $A = A + 10$ Write(A)




- Test for Conflict serializability
- Test for schedule equivalence
- Precedence Graph
  - $T_1$  executes write(A) before  $T_2$  executes read(A) (edge from  $T_1 \rightarrow T_2$ )
  - $T_2$  executes write(B) before  $T_1$  executes read(B) (edge from  $T_2 \rightarrow T_1$ )
  - So S5 is not conflict serializable
- Schedule equivalence (A+B)
  - Before transaction,  $A+B = 1500$
  - After transaction,  $A+B = 1500$
  - So  $S_5$  and  $S_5'$  (Any schedule equivalent to  $S_5$  by swapping non conflicting instructions ) are equivalent schedules
- It is possible to have two schedules that produce same outcome but are not conflict serializable
- Schedule equivalence have **less-stringent** definitions

# Recoverability

- If a transaction fails during its execution

Then, the partial executed failed transaction must be rolled back, thereby undoing all its effects as to preserve the Atomicity property.

$T_{10}$	$T_{11}$	$T_{12}$
read (A) read (B) write (A)  abort	 read (A) write (A)	  read (A)



## Recoverable schedules

In a concurrent transaction execution failure of transaction requires rolling back of that transaction along with those transaction, which are dependent on failed transaction in order to preserve atomicity.

# Recoverable Schedules

- **Recoverable schedule** — if a transaction  $T_j$  reads a data item previously written by a transaction  $T_i$ , then the commit operation of  $T_i$  **must** appear before the commit operation of  $T_j$ .

- To make this schedule recoverable,  $T_9$  have to delay committing until after  $T_8$  commits

$T_8$	$T_9$
read (A) write (A)	
	read (A) commit
read (B) abort	

- This is a **non recoverable** schedule because
  - If  $T_8$  fails before it commits, then  $T_9$  reads new value of A. i.e.  $T_9$  is dependent on  $T_8$ .
  - Therefore  $T_9$  should also be aborted along with  $T_8$
  - But  $T_9$  already committed with a inconsistent database state.



# Cascading Rollbacks

- A single transaction failure may lead to a series of **transaction rollbacks**.

- $T_{10}$  writes A, read by  $T_{11}$
- $T_{11}$  writes A, read by  $T_{12}$
- $T_{12}$  depends on  $T_{11}$  and  $T_{11}$  depends on  $T_{10}$

$T_{10}$	$T_{11}$	$T_{12}$
read (A) read (B) write (A)	read (A) write (A)	read (A)
abort		

- Now if  $T_{10}$  fails, then  $T_{11}$ , and  $T_{12}$  has also to be rolled back along with  $T_{10}$  due their interdependency.
- If a transaction failure leads to a series of rollbacks, is called **cascading rollback**
- It is undesirable as involves significant amount of work.

# Cascadeless Schedules

- It is desirable to restrict the schedules so that cascading rollback can't occur.
- These schedules are called **cascadeless** schedule.
- i.e. for each pair of transactions  $T_i$  and  $T_j$  such that  $T_j$  reads a data item previously written by  $T_i$ , the commit operation of  $T_i$  appears before the read operation of  $T_j$ .
- Every cascadeless schedule is also recoverable

End of Chapter  
Thank you