



# Concurrency Control

Instructor : Nitesh Kumar Jha

ITER,S'O'A(DEEMED TO BE  
UNIVERSITY)

Oct 2018

# Concurrency Control

- It is a mechanism to ensure isolation in a concurrent execution scenario
- Achieved using the concept of **mutual exclusion**
  - i.e. while one transaction is accessing a data item, no other transaction is allowed to modify that data item.
- Mutual exclusion is achieved using logical **locks**.
  - Locks are granted/revoked by a central **concurrency control manager**.
  - i.e. Transactions request it to grant a lock

# Locks

- Data items can be locked in two modes :
  - *Shared*: Data item can only be read. Requested using **lock-S** instruction. It can be shared with other transactions
  - *Exclusive*: Data item can be both read as well as written. It is requested using **lock-X** instruction. It can't be shared with other transactions

- Lock-compatibility matrix

- Shared locks may be granted to multiple transactions simultaneously.
- Exclusive locks can't be granted to multiple transactions simultaneously

	S	X
S	true	false
X	false	false

# Locks

- Let A and B are two accounts that are accessed by transactions  $T_1$  and  $T_2$ . Transaction  $T_1$  transfers \$50 from account B to account A. Transaction  $T_2$  displays the total amount of money in accounts A and B i.e., sum  $A+B$ .
- If these transaction are executed serially,  $T_1 T_2$  or  $T_2 T_1$ , It will be consistent.
- What happens if these are executed concurrently???

$T_1$	$T_2$
Lock-X(B)	Lock-S(A)
Read(B)	Read(A)
$B=B-50$	Unlock(A)
Write(B)	Lock-S(B)
Unlock(B)	Read(B)
Lock-X(A)	Unlock(B)
Read(A)	Display ( $A+B$ )
$A=A+50$	
Write(A)	
Unlock(A)	

# Locking Example

T <sub>1</sub>	T <sub>2</sub>	CC Manager	T <sub>1</sub>	T <sub>2</sub>	CC Manager
Lock-X(B)	A=500	Grant-X(B, T <sub>1</sub> )		Read(A)	Grant-S(A, T <sub>2</sub> )
	B=1000			Unlock(A)	1450
Read(B)	1000			Display (A+B)	
B=B-50	950				
Write(B)		Grant-S(B, T <sub>2</sub> )	Lock-X(A)		Grant-X(A, T <sub>1</sub> )
Unlock(B)			Read(A)		
	Lock-S(B)		A=A+50		
	Read(B)		Write(A)		
	Unlock(B)		Unlock(B)		
	Lock-S(A)				

- Produces inconsistent result
- Transactions must hold the lock on a data item till it access that item.
- It is not necessarily desirable for a transaction to unlock a data item immediately after its final access of that item, since the serializability may not be ensured.

# Delayed Unlocking

- Unlocking is delayed to the end of the transactions
- This schedule produces consistent result

$T_3$ : lock-X( $B$ );  
read( $B$ );  
 $B := B - 50$ ;  
write( $B$ );  
lock-X( $A$ );  
read( $A$ );  
 $A := A + 50$ ;  
write( $A$ );  
unlock( $B$ );  
unlock( $A$ ).

$T_4$ : lock-S( $A$ );  
read( $A$ );  
lock-S( $B$ );  
read( $B$ );  
display( $A + B$ );  
unlock( $A$ );  
unlock( $B$ ).

- But this technique may lead to an undesirable scenario called deadlock

# Deadlock due to Hold-and-Wait

$T_5$	$T_6$	CC Manager
Lock-X(B)		Grant-X(B, $T_1$ )
Read(B)		
$B = B - 50$		
Write(B)	Lock-S(A)	Grant-S(A, $T_2$ )
	Read(A)	
Lock-X(A)	Lock-S(B)	

Deadlock

- ❑ When a transaction ( $T_5$ ) delays unlocking on its locked data items (B) and requests to acquire a lock on new data items (A) that is already locked by another transaction ( $T_6$ )
- ❑ This is called a Hold-and-Wait situation

# Deadlock due to Hold-and-Wait

$T_5$	$T_6$	CC Manager
Lock-X(B)		Grant-X(B, $T_1$ )
Read(B)		
B=B-50		
Write(B)		
	Lock-S(A)	
		Grant-S(A, $T_2$ )
	Read(A)	
	Lock-S(B)	
Lock-X(A)		
Deadlock		

- Deadlock is a state where neither of these transactions can ever proceed with its normal execution. It can be resolved by forcibly rolling back one or more participating transactions.
- Lock based concurrency control needs the transaction to follow a set of rules called **locking protocol**
- A locking protocol is a set of rules followed by all transactions while requesting and releasing locks.
- Locking protocols restrict the set of possible schedules.



# The Two-Phase Locking Protocol

- This protocol ensures conflict-serializable schedules.
- It requires the transactions execute in two phases
- Phase 1: Growing Phase
  - Transaction may obtain locks
  - Transaction may not release locks
- Phase 2: Shrinking Phase
  - Transaction may release locks
  - Transaction may not obtain locks
- This protocol assures serializability.

# Lock Conversions

- Two-phase locking with lock conversions:
  - First Phase:
    - can acquire a lock-S on item
    - can acquire a lock-X on item
    - can convert a lock-S to a lock-X (upgrade)
  - Second Phase:
    - can release a lock-S
    - can release a lock-X
    - can convert a lock-X to a lock-S (downgrade)
- This protocol assures serializability. But still relies on the programmer to insert the various locking instructions.

# Automatic Acquisition of Locks

Not IMP

- A transaction  $T_i$  issues the standard read/write instruction, without explicit locking calls.
- The operation **read**( $D$ ) is processed as:
  - if**  $T_i$  has a lock on  $D$
  - then**
  - read( $D$ )
  - else begin**
  - if necessary wait until no other transaction has a **lock-X** on  $D$
  - grant  $T_i$  a **lock-S** on  $D$ ;
  - read( $D$ )
  - end**

# Automatic Acquisition of Locks (Cont.)

Not IMP

- **write**( $D$ ) is processed as:

if  $T_i$  has a **lock-X** on  $D$

then

write( $D$ )

else begin

if necessary wait until no other transaction has any lock on  $D$ ,

if  $T_i$  has a **lock-S** on  $D$

then

upgrade lock on  $D$  to **lock-X**

else

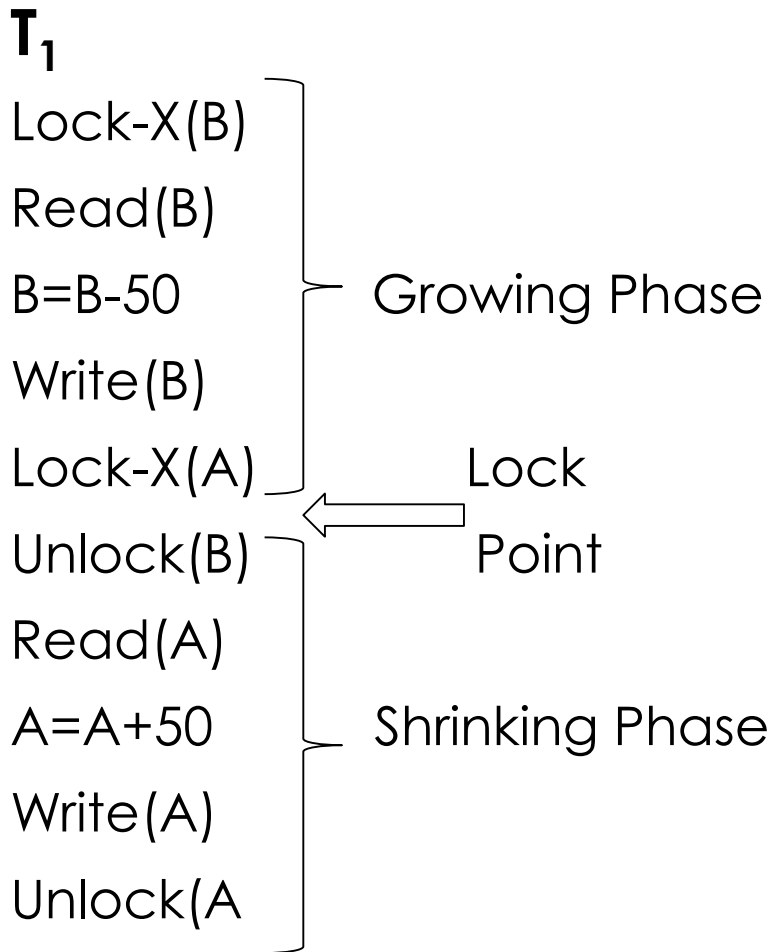
grant  $T_i$  a **lock-X** on  $D$

write( $D$ )

end;


- All locks are released after commit or abort

# The Two-Phase Locking Protocol



- ❑ The point in the schedule where transaction has obtained its final lock(end of its growing phase) is called **lock point** of the transaction.
- ❑ Now transactions can be ordered according to their lock points.

# Two Phase Locking Example

T <sub>1</sub>	T <sub>2</sub>	CC Manager	T <sub>1</sub>	T <sub>2</sub>	CC Manager
Lock-X(B)		Grant-X(B, T <sub>1</sub> )		Read(A)	
Read(B)	<b>Not proceed further</b>			Display (A+B)	1450
B=B-50	<b>So T<sub>1</sub> Will proceed first then T<sub>2</sub></b>			Unlock(B)	
Write(B)				Unlock(A)	
	Lock-S(B)	Grant-S(B, T <sub>2</sub> )	Lock-X(A)		Grant-X(A, T1)
	Read(B)		Read(A)		
	Lock-S(A)	Grant-S(A, T <sub>2</sub> )	A=A+50		
			Write(A)		
			Unlock(B)		
			Unlock(A)		

- Produces consistent result as transaction will be serial T<sub>1</sub> -> T<sub>2</sub>
- Transactions holds the lock on a data item till the last.

# Properties of 2PL

- 2PL ensures conflict serializability. The contributing transactions are isolated w.r.t. the lock point. (point at which growing phase ends and shrinking phase starts)
- It does not ensure deadlock free execution. In the event of deadlock participating transactions are rolled back. Consider

Previous Example here ----->

T <sub>5</sub>	T <sub>6</sub>	CC Manager
Lock-X(B)		Grant-X(B, T <sub>1</sub> )
Read(B)		
B=B-50		
Write(B)		
	Lock-S(A)	
	Read(A)	Grant-S(A, T <sub>2</sub> )
	Lock-S(B)	
Lock-X(A)		
Deadlock		

- It ensures recoverability but does not safeguard against cascading rollback.

# Deadlocks

- Consider the partial schedule

$T_3$	$T_4$
lock-x ( $B$ )	
read ( $B$ )	
$B := B - 50$	
write ( $B$ )	
	lock-s ( $A$ )
	read ( $A$ )
	lock-s ( $B$ )
lock-x ( $A$ )	

- Neither  $T_3$  nor  $T_4$  can make progress — executing **lock-S( $B$ )** causes  $T_4$  to wait for  $T_3$  to release its lock on  $B$ , while executing **lock-X( $A$ )** causes  $T_3$  to wait for  $T_4$  to release its lock on  $A$ .
- Such a situation is called a **deadlock**.
  - To handle a deadlock one of  $T_3$  or  $T_4$  must be rolled back and its locks released.



# Deadlocks (Cont.)

- The potential for deadlock exists in most locking protocols. Deadlocks are a necessary evil.
- When a deadlock occurs there is a possibility of cascading roll-backs.
- Cascading roll-back is possible under two-phase locking. To avoid this, we must follow a modified protocol called **strict two-phase locking** -- a transaction must hold all its exclusive locks till it commits/aborts.

# Strict two Phase locking Protocol

- **Strict two phase locking** is an enhanced 2PL that ensures cascadeless recovery.
- Strict 2PL demands that not only the locking and unlocking be in two phases, all the exclusive mode locks must be held by the transaction till the transaction commits.
- This requirement ensured that any data written by uncommitted transactions are locked in Exclusive mode until the transaction commits, preventing any other transaction from reading the data.
- **Rigorous two-phase locking** is even stricter. Here, *all* locks are held till commit/abort. In this protocol transactions can be serialized in the order in which they commit.

# Deadlocks (Cont.)

- Strict Two-phase locking *does not* ensure freedom from deadlocks.
- In addition to deadlocks, there is a possibility of **starvation**.
- **Starvation** occurs if the concurrency control manager is badly designed. For example:
  - A transaction may be waiting for an X-lock on an item, while a sequence of other transactions request and are granted an S-lock on the same item.
  - The same transaction is repeatedly rolled back due to deadlocks.
- Concurrency control manager can be designed to prevent starvation.


# Implementation of Locking

Not IMP

- A **lock manager** can be implemented as a separate process to which transactions send lock and unlock requests
- The lock manager replies to a lock request by sending a lock grant messages (or a message asking the transaction to roll back, in case of a deadlock)
- The requesting transaction waits until its request is answered
- The lock manager maintains a data-structure called a **lock table** to record granted locks and pending requests
- The lock table is usually implemented as an in-memory hash table indexed on the name of the data item being locked

# Lock Table

Not IMP



- Dark blue rectangles indicate granted locks; light blue indicate waiting requests
- Lock table also records the type of lock granted or requested
- New request is added to the end of the queue of requests for the data item, and granted if it is compatible with all earlier locks
- Unlock requests result in the request being deleted, and later requests are checked to see if they can now be granted
- If transaction aborts, all waiting or granted requests of the transaction are deleted
  - lock manager may keep a list of locks held by each transaction, to implement this efficiently

# Deadlock Handling

Not IMP

- System is deadlocked if there is a set of transactions such that every transaction in the set is waiting for another transaction in the set.
- **Deadlock prevention** protocols ensure that the system will *never* enter into a deadlock state.  
Some prevention strategies :
  - Require that each transaction locks all its data items before it begins execution (predeclaration).
  - Impose partial ordering of all data items and require that a transaction can lock data items only in the order specified by the partial order.

# More Deadlock Prevention Strategies

- Following schemes use transaction timestamps for the sake of deadlock prevention alone. Not IMP
- **wait-die** scheme — non-preemptive
  - older transaction may wait for younger one to release data item. (older means smaller timestamp) Younger transactions never wait for older ones; they are rolled back instead.
  - a transaction may die several times before acquiring needed data item
- **wound-wait** scheme — preemptive
  - If a younger transaction holds a data item requested by older transaction, the younger transaction is aborted and rolled back. Younger transactions may wait for older ones to release a data item.
  - may be fewer rollbacks than *wait-die* scheme.

# Deadlock prevention (Cont.)

Not IMP

- Both in *wait-die* and in *wound-wait* schemes, a rolled back transactions is restarted with its original timestamp. Older transactions thus have precedence over newer ones, and starvation is hence avoided.
- **Timeout-Based Schemes:**
  - a transaction waits for a lock only for a specified amount of time. If the lock has not been granted within that time, the transaction is rolled back and restarted,
  - Thus, deadlocks are not possible
  - simple to implement; but starvation is possible. Also difficult to determine good value of the timeout interval.




# Deadlock Detection

Not IMP


- Deadlocks can be described as a *wait-for graph*, which consists of a pair  $G = (V, E)$ ,
  - $V$  is a set of vertices (all the transactions in the system)
  - $E$  is a set of edges; each element is an ordered pair  $T_i \rightarrow T_j$ .
- If  $T_i \rightarrow T_j$  is in  $E$ , then there is a directed edge from  $T_i$  to  $T_j$ , implying that  $T_i$  is waiting for  $T_j$  to release a data item.
- When  $T_i$  requests a data item currently being held by  $T_j$ , then the edge  $T_i \rightarrow T_j$  is inserted in the wait-for graph. This edge is removed only when  $T_j$  is no longer holding a data item needed by  $T_i$ .
- The system is in a deadlock state if and only if the wait-for graph has a cycle. Must invoke a deadlock-detection algorithm periodically to look for cycles.

# Deadlock Detection (Cont.)

Not IMP



Wait-for graph without a cycle



Wait-for graph with a cycle

# Deadlock Recovery

Not IMP

- When deadlock is detected :
  - Some transaction will have to rolled back (made a victim) to break deadlock. Select that transaction as victim that will incur minimum cost.
  - Rollback -- determine how far to roll back transaction
    - ▶ **Total rollback**: Abort the transaction and then restart it.
    - ▶ More effective to roll back transaction only as far as necessary to break deadlock.
  - Starvation happens if same transaction is always chosen as victim. Include the number of rollbacks in the cost factor to avoid starvation

# Timestamp-Based Protocols

- Each transaction is issued a timestamp when it enters the system. If an old transaction  $T_i$  has time-stamp  $TS(T_i)$ , a new transaction  $T_j$  is assigned time-stamp  $TS(T_j)$  such that  $TS(T_i) < TS(T_j)$ .
- The protocol manages concurrent execution such that the time-stamps determine the serializability order.
- In order to assure such behavior, the protocol maintains for each data  $Q$  two timestamp values:
  - **W-timestamp**( $Q$ ) is the largest time-stamp of any transaction that executed **write**( $Q$ ) successfully.
  - **R-timestamp**( $Q$ ) is the largest time-stamp of any transaction that executed **read**( $Q$ ) successfully.

# Timestamp-Based Protocols (Cont.)


- The timestamp ordering protocol ensures that any conflicting **read** and **write** operations are executed in timestamp order.
- Suppose a transaction  $T_i$  issues a **read**(Q)
  1. If  $TS(T_i) \leq \mathbf{W}\text{-timestamp}(Q)$ , then  $T_i$  needs to read a value of Q that was already overwritten by a younger transaction.
    - Hence, the **read** operation is rejected, and  $T_i$  is rolled back.
  2. If  $TS(T_i) \geq \mathbf{W}\text{-timestamp}(Q)$ , then the **read** operation is executed, and  $\mathbf{R}\text{-timestamp}(Q)$  is set to  $\mathbf{max}(\mathbf{R}\text{-timestamp}(Q), TS(T_i))$ .

# Timestamp-Based Protocols (Cont.)

- Suppose that transaction  $T_i$  issues **write**(Q).
  1. If  $TS(T_i) < R\text{-timestamp}(Q)$ , then the value of Q that  $T_i$  is producing was needed previously, and the system assumed that that value would never be produced.
    - Hence, the **write** operation is rejected, and  $T_i$  is rolled back.
  2. If  $TS(T_i) < W\text{-timestamp}(Q)$ , then  $T_i$  is attempting to write an obsolete value of Q.
    - Hence, this **write** operation is rejected, and  $T_i$  is rolled back.
  3. Otherwise, the **write** operation is executed, and  $W\text{-timestamp}(Q)$  is set to  $TS(T_i)$ .

# Correctness of Timestamp-Ordering Protocol

- The timestamp-ordering protocol guarantees serializability since all the arcs in the precedence graph are of the form:



Thus, there will be no cycles in the precedence graph

- Timestamp protocol ensures freedom from deadlock as no transaction ever waits.
- But the schedule may not be cascade-free, and may not even be recoverable.

# Recoverability and Cascade Freedom

- Problem with timestamp-ordering protocol:
  - Suppose  $T_i$  aborts, but  $T_j$  has read a data item written by  $T_i$
  - Then  $T_j$  must abort; if  $T_j$  had been allowed to commit earlier, the schedule is not recoverable.
  - Further, any transaction that has read a data item written by  $T_j$  must abort
  - This can lead to cascading rollback --- that is, a chain of rollbacks
- Solution 1:
  - A transaction is structured such that its writes are all performed at the end of its processing
  - All writes of a transaction form an atomic action; no transaction may execute while a transaction is being written
  - A transaction that aborts is restarted with a new timestamp
- Solution 2: Limited form of locking: wait for data to be committed before reading it
- Solution 3: Use commit dependencies to ensure recoverability



# Thomas' Write Rule

Not IMP

- Modified version of the timestamp-ordering protocol in which obsolete **write** operations may be ignored under certain circumstances.
- When  $T_i$  attempts to write data item  $Q$ , if  $TS(T_i) < W\text{-timestamp}(Q)$ , then  $T_i$  is attempting to write an obsolete value of  $\{Q\}$ .
  - Rather than rolling back  $T_i$  as the timestamp ordering protocol would have done, this **{write}** operation can be ignored.
- Otherwise this protocol is the same as the timestamp ordering protocol.
- Thomas' Write Rule allows greater potential concurrency.
  - Allows some view-serializable schedules that are not conflict-serializable.

# Validation-Based Protocol

Not IMP

- Execution of transaction  $T_i$  is done in three phases.
  1. **Read and execution phase:** Transaction  $T_i$  writes only to temporary local variables
  2. **Validation phase:** Transaction  $T_i$  performs a "validation test" to determine if local variables can be written without violating serializability.
  3. **Write phase:** If  $T_i$  is validated, the updates are applied to the database; otherwise,  $T_i$  is rolled back.
- The three phases of concurrently executing transactions can be interleaved, but each transaction must go through the three phases in that order.
  - Assume for simplicity that the validation and write phase occur together, atomically and serially
    - ▶ I.e., only one transaction executes validation/write at a time.
- Also called as **optimistic concurrency control** since transaction executes fully in the hope that all will go well during validation

# Schedule Produced by Validation

Not IMP

- Example of schedule produced using validation

$T_{25}$	$T_{26}$
read (B)	read (B) $B := B - 50$ read (A) $A := A + 50$
read (A) $\langle \text{validate} \rangle$ display (A + B)	$\langle \text{validate} \rangle$ write (B) write (A)

# Multiversion Schemes

Not IMP

- Multiversion schemes keep old versions of data item to increase concurrency.
  - Multiversion Timestamp Ordering
  - Multiversion Two-Phase Locking
- Each successful **write** results in the creation of a new version of the data item written.
- Use timestamps to label versions.
- When a **read**(Q) operation is issued, select an appropriate version of Q based on the timestamp of the transaction, and return the value of the selected version.
- **reads** never have to wait as an appropriate version is returned immediately.

# Multiversion Timestamp Ordering

Not IMP

- Each data item  $Q$  has a sequence of versions  $\langle Q_1, Q_2, \dots, Q_m \rangle$ . Each version  $Q_k$  contains three data fields:
  - **Content** -- the value of version  $Q_k$ .
  - **W-timestamp**( $Q_k$ ) -- timestamp of the transaction that created (wrote) version  $Q_k$
  - **R-timestamp**( $Q_k$ ) -- largest timestamp of a transaction that successfully read version  $Q_k$
- When a transaction  $T_i$  creates a new version  $Q_k$  of  $Q$ ,  $Q_k$ 's W-timestamp and R-timestamp are initialized to  $TS(T_i)$ .
- R-timestamp of  $Q_k$  is updated whenever a transaction  $T_j$  reads  $Q_k$ , and  $TS(T_j) > \text{R-timestamp}(Q_k)$ .

# Multiversion Timestamp Ordering (Cont)

Not IMP

- Suppose that transaction  $T_i$  issues a **read**(Q) or **write**(Q) operation. Let  $Q_k$  denote the version of Q whose write timestamp is the largest write timestamp less than or equal to  $TS(T_i)$ .
  1. If transaction  $T_i$  issues a **read**(Q), then the value returned is the content of version  $Q_k$ .
  2. If transaction  $T_i$  issues a **write**(Q)
    1. if  $TS(T_i) < \text{R-timestamp}(Q_k)$ , then transaction  $T_i$  is rolled back.
    2. if  $TS(T_i) = \text{W-timestamp}(Q_k)$ , the contents of  $Q_k$  are overwritten
    3. else a new version of Q is created.
- Observe that
  - Reads always succeed
  - A write by  $T_i$  is rejected if some other transaction  $T_j$  that (in the serialization order defined by the timestamp values) should read  $T_i$ 's write, has already read a version created by a transaction older than  $T_i$ .
- Protocol guarantees serializability

# MVCC: Implementation Issues

Not IMP

- Creation of multiple versions increases storage overhead
  - Extra tuples
  - Extra space in each tuple for storing version information
- Versions can, however, be garbage collected
  - E.g. if Q has two versions Q5 and Q9, and the oldest active transaction has timestamp  $> 9$ , then Q5 will never be required again

**End of Lecture**  
**Thank You**